



Test Foxy Audit yourself

Foxy Audit proves an AI system didn't leak or tamper with sensitive data — **without ever seeing that data**. Verify that claim yourself in a few minutes.

1 Your evaluation accounts

Five ready-to-use accounts, each with **20,000 audit-event credits**. Log in at app.foxyaudit.tech — the dashboard opens straight away.

#	EMAIL	PASSWORD	CREDITS
1	judge1@foxyaudit.tech	Foxy-Judge-2026	20,000
2	judge2@foxyaudit.tech	Foxy-Judge-2026	20,000
3	judge3@foxyaudit.tech	Foxy-Judge-2026	20,000
4	judge4@foxyaudit.tech	Foxy-Judge-2026	20,000
5	judge5@foxyaudit.tech	Foxy-Judge-2026	20,000

Your API key: after logging in, open **Access** in the left sidebar and copy (or create) an API key. It's shown once — paste it into the test script below.

2 Two-minute proof — no account needed

From the repository root, with **Python 3.9+**:

```
# install the SDK (also on PyPI)
pip install foxy-audit

# build a real hash-chain, verify it, then watch tamper-detection fire
python demo/offline_demo.py

# drive the host-side guard with a mock LLM - PHI/PII/secrets/injection blocked locally
python demo/mock_llm.py --scenario all
```

Neither command contacts our servers or needs a key — they run the *real* commitment, chain, and policy logic against synthetic input.

3 Full test — against your live account

```
# 1. log in (table above) → Access → copy your API key, then:
setx FOXY_API_KEY "foxy_sk_... your key ..."      # new PowerShell after this
python demo/run_demo.py                             # captures events to YOUR ledger
```

Then, back in the dashboard at app.foxyaudit.tech:

Ledger

See the allowed, **blocked**, and **redacted** events you just produced — hashes only, never raw text.

Policy → AI Judge

Pick who grades your events — **GPT-5.6**, Gemini, or your own key. Content-blind either way.

Export → Compliance Passport

One click → a signed PDF: what was captured, what was blocked, chain intact.

Verify

Export the ledger and run `python verifier/foxy_verify.py` — change one byte and it flags the exact row.

4 See GPT-5.6 grade (optional)

On **Policy** → “**AI Judge — provider & keys**”, choose **OpenAI**, paste your OpenAI key, and Save. Your events are then graded by **GPT-5.6 via the OpenAI Responses API** — which receives only

hashes, token counts, and policy tags, **never the prompt or response**.

5 What you're verifying

- ✓ Sensitive prompts are **blocked on the host**, before the model is ever called.
- ✓ Only a **64-character hash** reaches the backend — the raw prompt/response never leaves your machine.
- ✓ The audit trail is **tamper-evident**: one changed byte breaks the chain from that row on.
- ✓ The AI judge (GPT-5.6) is **content-blind** — it grades metadata, not conversations.

Honest scope. Foxy Audit is an evidence and integrity layer — not a compliance certification, a network firewall, or proof that an un-instrumented model call was captured. Everything shown here is independently verifiable in the source code and the git history. · Evaluation accounts are throwaway and rate-limited.